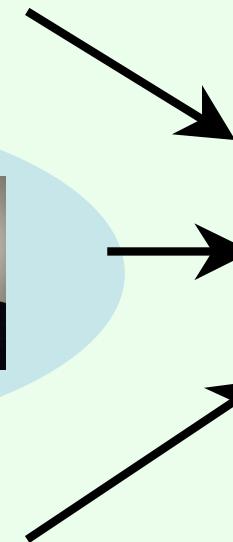
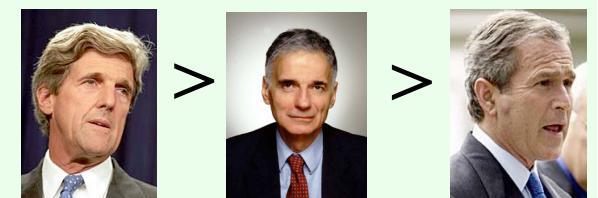
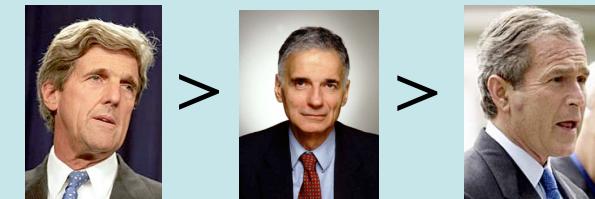


Eliciting Single-Peaked Preferences Using Comparison Queries

Vincent Conitzer
conitzer@cs.duke.edu

Voting



Pairwise elections



>

>



>



>



two votes prefer Kerry to Bush



>



two votes prefer Kerry to Nader



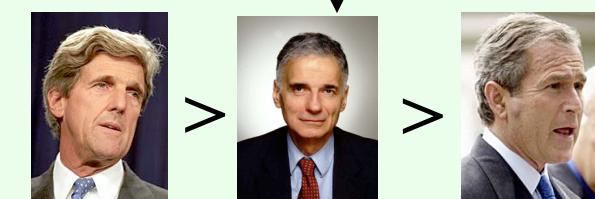
>



two votes prefer Nader to Bush



>



Condorcet cycles

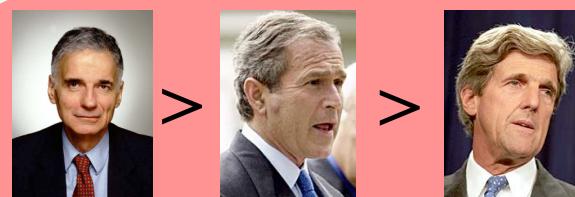
two votes prefer Bush to Kerry



two votes prefer Kerry to Nader



two votes prefer Nader to Bush



“weird” preferences

Single-peaked preferences [Black 48]

- Suppose alternatives are ordered on a line from left to right (the alternatives' **positions**)
- E.g. d - b - e - f - a - c
 - Left-wing vs. right-wing political candidates
 - Perhaps the alternatives are numbers, e.g. voting over the size of the budget
 - Voting over locations along a road
 - ...
- An agent's preferences are **single-peaked** with respect to these positions if the agent prefers alternatives that lie closer to her most preferred alternative (on each side)
- $f > e > a > d > c > b$ is **not** single peaked with respect to above positions: d is ranked above b, but b is closer to f and on the same side as d
- $f > e > b > a > c > d$ **is** single-peaked

Nice properties of single-peaked preferences

- Suppose every voter's preferences are single-peaked (with respect to the same positions for the alternatives)
- If a wins the pairwise election between a and b,
- and b wins the pairwise election between b and c,
- then a must win the pairwise election between a and c
 - I.e. no Condorcet cycles
- So we can use pairwise elections to determine the ranking
- This is also strategy-proof
- (**Gibbard-Satterthwaite theorem**: for general preferences, no reasonable deterministic voting rule is strategy-proof)

Preference elicitation

- Direct mechanisms ask each agent to reveal complete preferences
 - In voting, each agent gives an entire ranking
- Can be cumbersome to agents
 - Have to decide and communicate entire preferences without any help
 - Especially hard if there are many alternatives
- In preference elicitation, the center (elicitor) repeatedly asks agents “natural” queries about their preferences
 - E.g. comparison queries: do you prefer a to b?
- In this paper, the elicitor wants to learn each agent’s complete preferences, using comparison queries
- How many queries are needed?

Eliciting general preferences (not single-peaked)

- Discover the full ranking $>$ of the m alternatives based on comparison queries
- Equivalent to sorting a list of m elements using only binary comparisons
- E.g. MergeSort algorithm solves this with $O(m \log m)$ queries
- Any algorithm is $\Omega(m \log m)$
- With n voters, many voting rules require $\Omega(nm \log m)$ communication even just to determine the winner
[Conitzer & Sandholm EC05]

Eliciting preferences given positions

- Voter's preferences: $b > c > e > f > a > d$ (unknown)
- Positions: $e - c - b - f - a - d$ (known)
- Let us find the most preferred alternative first
- “ $b > f?$ ” “Yes”
 - Tells us that most preferred alternative must be e, c, b
 - \sim binary search
- “ $c > b?$ ” “No”
 - So b must be most preferred
 - Next-ranked alternative must be c or f
- “ $c > f?$ ” “Yes”
 - Next-ranked alternative must be e or f
- “ $e > f?$ ” “Yes”
 - Now we know the ranking must be $b > c > e > f > a > d$

How many queries does this take?

- Finding the most preferred alternative takes at most $1 + \log m$ queries
 - Binary search
- The remainder will require at most $m - 2$ queries
 - Each query allows us to add the next alternative to the ranking

What if we do not know the positions?

- Any preferences are single-peaked with respect to some positions
- E.g. $f > e > a > d > c > b$ is consistent with respect to
 - f - e - a - d - c - b
 - d - e - f - a - c - b
 - many other positionings
- So eliciting the first voter's preferences will require $\Omega(m \log m)$ queries
- Once we know one voter's preferences, we know something about the positions
- Will show that this is enough information to need only $O(m)$ queries for next voter

Eliciting preferences using another voter's preferences (stage 1)

- Positions: $e - c - b - f - a - d$ (unknown)
- Current voter's preferences: $c > e > b > f > a > d$ (unknown)
- Previous voter's preferences: $a > d > f > b > c > e$ (known)
- Let us find the most preferred alternative first
- Cannot use binary search this time, just do one at a time
- “ $a > d?$ ” “Yes”
- “ $a > f?$ ” “No”
- “ $f > b?$ ” “No”
- “ $b > c?$ ” “No”
- “ $c > e?$ ” “Yes”
- So most preferred alternative is c

Eliciting preferences using another voter's preferences (stage 2)

- Positions: $e - c - b - f - a - d$ (unknown)
- Current voter's preferences: $c > e > b > f > a > d$ (unknown)
- Previous voter's preferences: $a > d > f > b > c > e$ (known)
- Let us find out which alternatives lie between a (previous voter's most preferred) and c (current voter's most preferred) in the positions
- Previous voter must prefer such alternatives to c
 - Could be d, f, b
- Current voter must prefer such alternatives to a
 - “ $d > a?$ ” “No”
 - “ $f > a?$ ” “Yes”
 - “ $b > a?$ ” “Yes”
- So b and f lie between a and c
- Current voter's preferences over a, c, b, f must be opposite of previous voter's, i.e. $c > b > f > a$

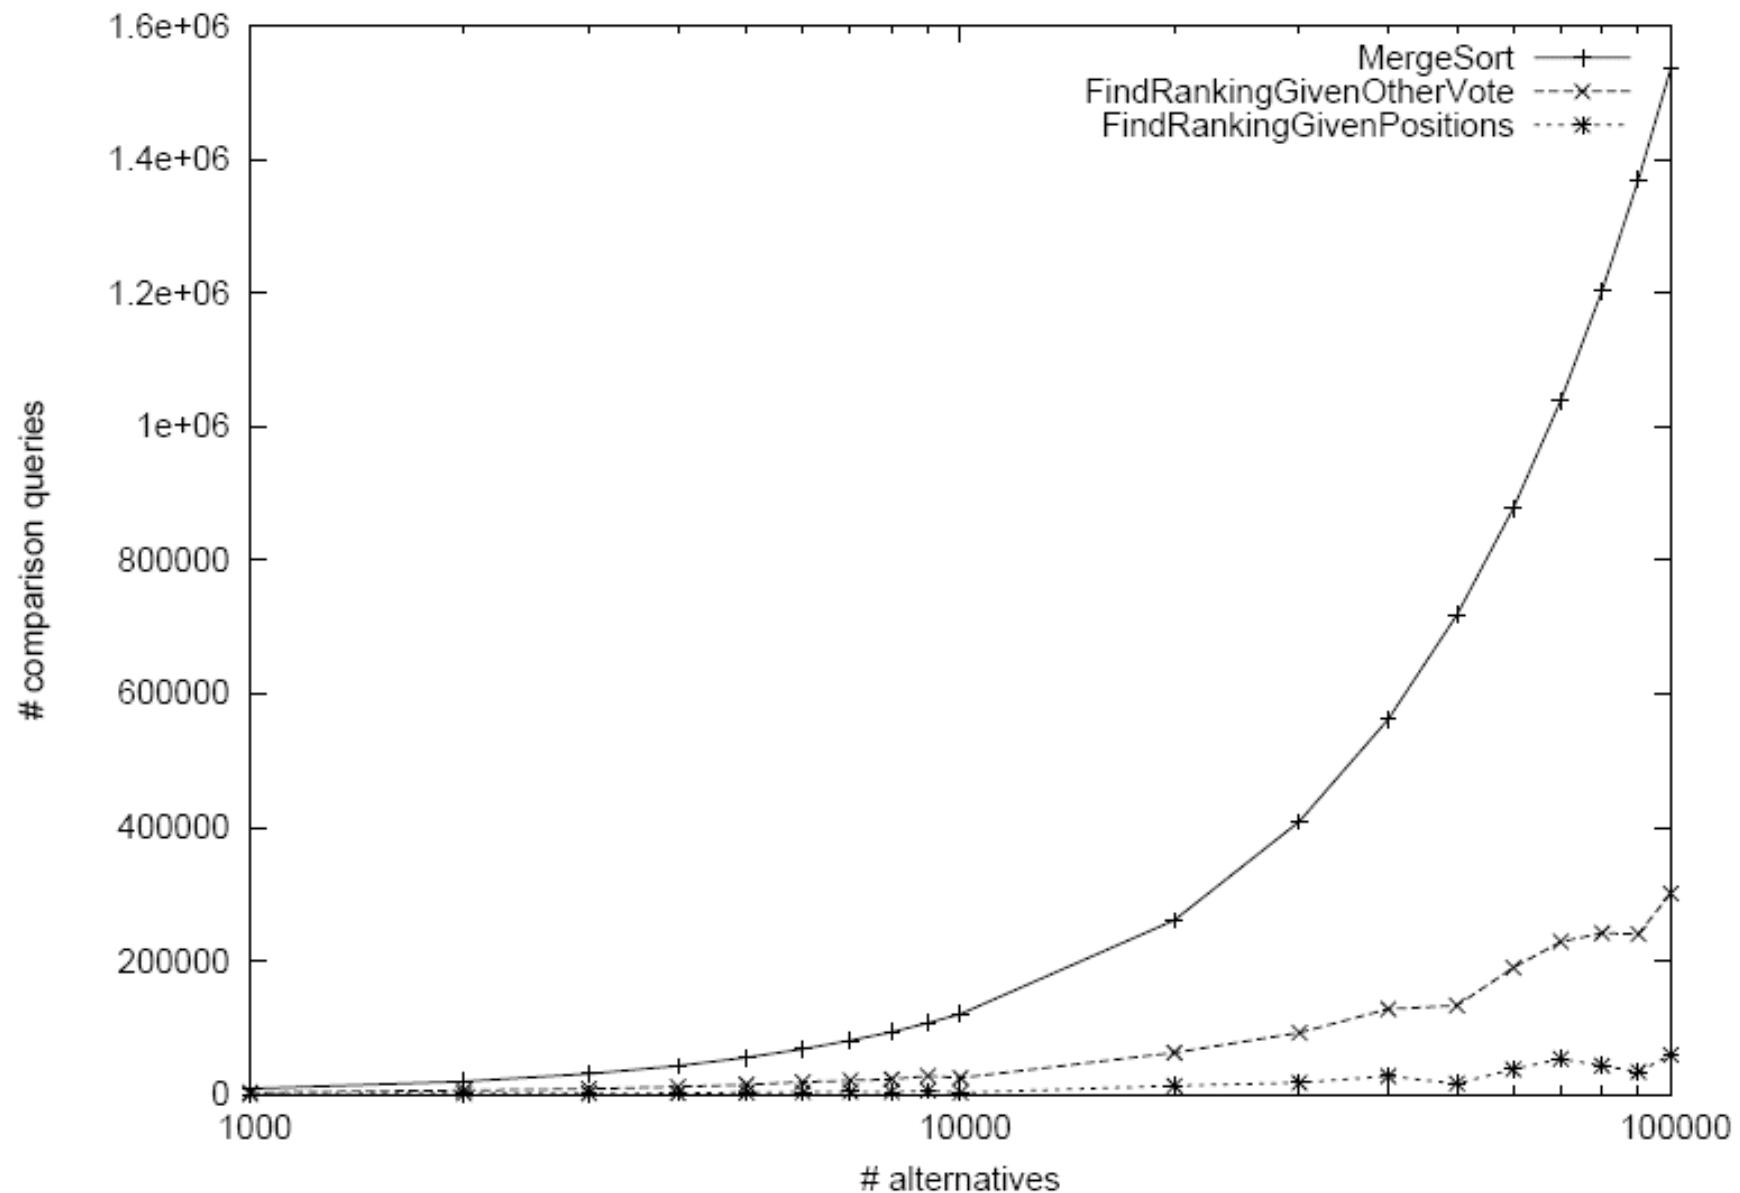
Eliciting preferences using another voter's preferences (stage 3)

- Positions: $e - c - b - f - a - d$ (unknown)
- Current voter's preferences: $c > e > b > f > a > d$ (unknown)
- Previous voter's preferences: $a > d > f > b > c > e$ (known)
- We know $c > b > f > a$; must integrate d and e
 - In order of previous voter's preferences, i.e. d before e
- Start by comparing to currently last-ranked alternative
- “ $d > a?$ ” “No”
 - Now we know $c > b > f > a > d$
- “ $e > d?$ ” “Yes”
 - e must lie on **opposite side** from d in positions, since known and current voters disagree on ranking of e and d
 - Start from the top...
- “ $e > b?$ ” “Yes”
 - Now we know $c > e > b > f > a > d$

How many queries does this take?

- Finding the most preferred alternative (stage 1) takes at most $m - 1$ queries
- Finding the alternatives between the previous and current voter's most preferred alternatives (stage 2) takes at most $m - 2$ queries
- Integrating the remaining alternatives (stage 3) requires at most $2m - 3$ queries
 - More complex argument
 - Requires keeping track of the worst-ranked alternative above which we will never insert another alternative
- Total upper bound is $4m - 6$ queries

Experimental results



Conclusions

- Determining general preferences requires $\Omega(m \log m)$ comparison queries
- If preferences are single-peaked and
 - the positions of the alternatives are known, or
 - at least one other voter's preferences are known,
- then preferences can be elicited using $O(m)$ queries
 - There is also an $\Omega(m)$ lower bound
- What about more general families of preferences?
 - E.g. alternatives take positions on the plane rather than the line
 - Many of the nice properties go away...

Thank you for your attention!